

# Software development

## Section 3: Neural Networks

Ioannis Emiris and Emmanouil Christoforou

Dept Informatics & Telecoms, NKUA

December 15, 2019

# Introduction


## Machine Learning

Machine learning (ML) uses statistical techniques to give computer systems the ability to “learn” (i.e. progressively improve performance on a specific task) from data, without being explicitly programmed.


## Neural Networks

Neural Networks (and Deep Learning) is a class of ML algorithms that use a cascade of layers of nonlinear processing units for feature extraction and transformation, mimicking the human neurons. Each successive layer uses the output from the previous layer as input.

# Live Neural Network





# Artificial Neural Network




# Neural Networks

- NN is a cascade of **hidden layers**
- Each layer is a **data transformation** function
  - weights/parameters determine the transformation
  - transformations are differentiable for improving output



# NN Layers of representations



- Input: Raw data
- Data transformed so that **irrelevant** information is **filtered out**

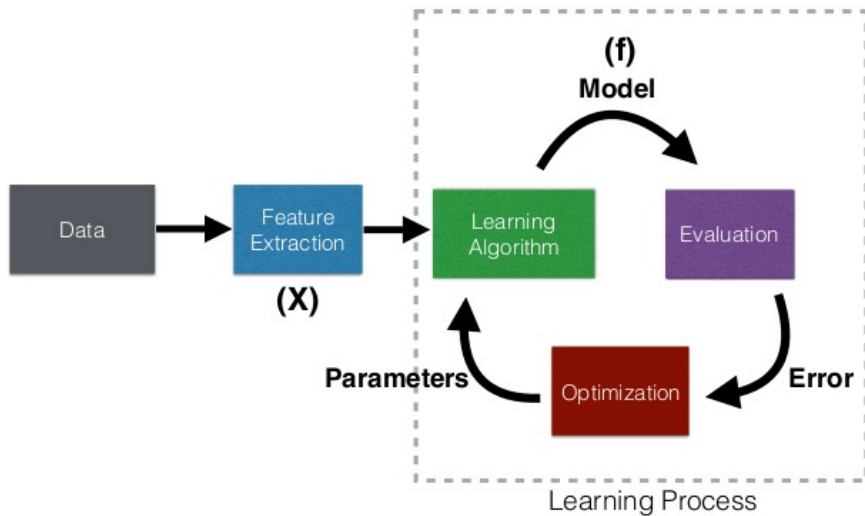
# Outline

1 Architecture

2 Layers

3 Usage

# Learning process




Predicted  $f(X)$ : Find  $f$  s.t.  $f(x)$  is close to **True value**




# Goal of training

**Goal of training** → find **weights** that minimise loss function





# Activation functions



- **Non-linear activation** transformation on (weighted) input
- Generate **non-linear mappings** from inputs to outputs
- Decide whether this neuron is **active** or not


# Activation functions (I)

## Traditional Non-Linear Activation Functions




## Modern Non-Linear Activation Functions

### Rectified Linear Unit (ReLU)




### Leaky ReLU



$\alpha = \text{small const. (e.g. 0.1)}$

### Exponential LU



# Loss functions based on problem

$y$  = true output label/value,  $\hat{y}$  = predicted class/value,  
 $n$  = number of training instances,  $M = \# \text{classes}$ .

- Regression

Mean Squared Error (MSE): 
$$\frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

Mean Absolute Error (MAE): 
$$\frac{1}{n} \sum_{i=1}^n |y_i - \hat{y}_i|$$


- Classification

binary cross-entropy: 
$$-\frac{1}{n} \sum_{i=1}^n (y_i \log \hat{y}_i + (1 - y_i) \log(1 - \hat{y}_i))$$


(multiclass, one label) cross-entropy: 
$$-\frac{1}{n} \sum_{i=1}^n \sum_j^M y_{ij} \log \hat{y}_{ij}$$

# Backward propagation

- Forward propagation: compute activations (black arrows)
- Backprop: compute derivatives optimizing parameters/weights (red)
- Gradient descent exploits chain rule  $\frac{dL}{dx} = \frac{dL}{dw} \frac{dw}{dx}$
- . . . to get **partial derivative of loss wrt weights**
- Learning rate: step to change weights




# Overfitting



Underfitting



Desired



Overfitting

# NN as universal approximators

## Universal Approximation Theorem


Let  $\xi$  be a non-constant, bounded, monotonically-increasing continuous activation function,  $f : [0, 1]^d \rightarrow \mathbb{R}$  continuous, and  $\epsilon > 0$ . Then,  $\exists n$ , parameters  $a, b \in \mathbb{R}^n$ ,  $W \in \mathbb{R}^{n \times d}$  s.t.

$$\left| \sum_{i=1}^n a_i \xi(w_i^T x + b_i) - f(x) \right| < \epsilon.$$

- Any  $f$  approximated arbitrarily well by NN with one hidden layer
- How many neurons? How to find the parameters?
- Does it generalize well? Does it overfit?
- Shallow nets work poorly for high-dimensional data like images

# Better Generalization with Greater Depth

- Empirical results show that **deeper** networks **generalize** better
- Test accuracy consistently increases when increasing depth



Transcribe multidigit numbers from photographs of addresses (Goodfellow'14)




# Outline

1 Architecture

2 Layers

3 Usage


# Fully Connected Layer



Matrix multiplication:  $W$  is a  $10 \times 3072$  real matrix.

# Convolution definition


- Pointwise matrix multiply
- The kernel (convolution mask) is a small matrix (or flipped)



# Convolution


Motivation: Sparse / Parameter sharing / Equivariant representations

Original:



Sharpen:

0	-1	0
-1	5	-1
0	-1	0



Blur:  $\frac{1}{9}$

1	1	1
1	1	1
1	1	1



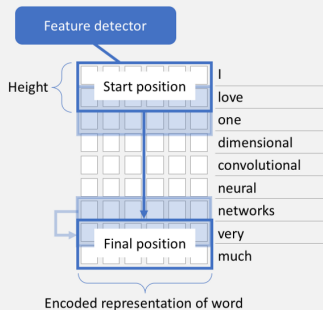
Edge  
detect:

-1	-1	-1
-1	8	-1
-1	-1	-1



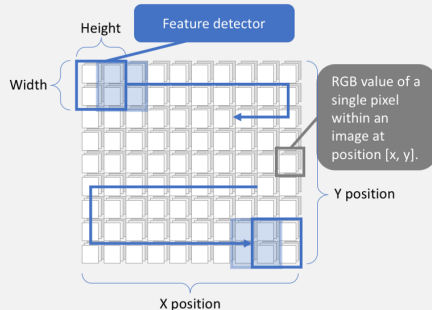
# 1D vs 2D convolution

## 1D Convolutional - Example



In this example for natural language processing, a sentence is made up of 9 words. Each word is a vector that represents a word as a low dimensional representation. The feature detector will always cover the whole word. The height determines how many words are considered when training the feature detector. In our example, the height is two. In this example the feature detector will iterate through the data 8 times.


## 2D Convolutional - Example



In this example for computer vision, each pixel within the image is represented by its x- and y position as well as three values (RGB). The feature detector has a dimension of 2 x 2 in our example. The feature detector will now slide both horizontally and vertically across the image.

# Pooling (Downsampling)

- Makes representations smaller and more manageable
- Helps make representation roughly invariant to small input translations
- **Max**, Avg,  $L^2$  norm, weighted avg distance from the central pixel





# Recurrent Neural Networks (RNN)

NNs mainly distinguished in:

- Feedforward neural networks (FNN):  
Features processed independently
- Recurrent neural networks (RNN):  
Features processed sequentially e.g. time series


RNN allow **cyclic** connections

→ fit best to process **sequential** data




# RNN input

## Feed-Forward Network Data




## Recurrent Network Data





# RNN: Type of sequences



e.g. **Video classification on frame level**

1-many: image captioning (image  $\rightarrow$  word sequence)


many-1: sentiment classification (word sequence  $\rightarrow$  sentiment)

many-many: machine translation (word sequence  $\rightarrow$  word seq.)

many-many: video classification on frame level

# Recurrent neural networks

- allow cyclic connections
- . . . essentially offer internal state (memory)
- keep track of arbitrarily long-term dependencies (boon/issue)
- internal state processes input  $x_t$ 's by applying recurrence formula at every time step  $t$ : new state  $h(t) \leftarrow f_w(h_{t-1}, x_t)$ .



Numerics: Back-propagated gradients may vanish ( $\rightarrow 0$ ) or explode ( $\rightarrow \infty$ ).

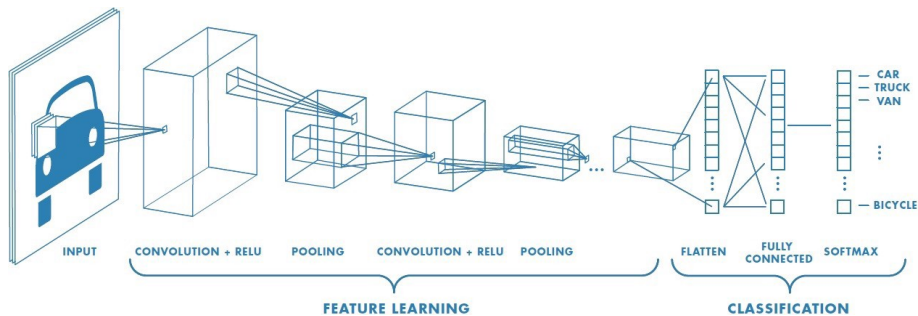
# Outline

1 Architecture

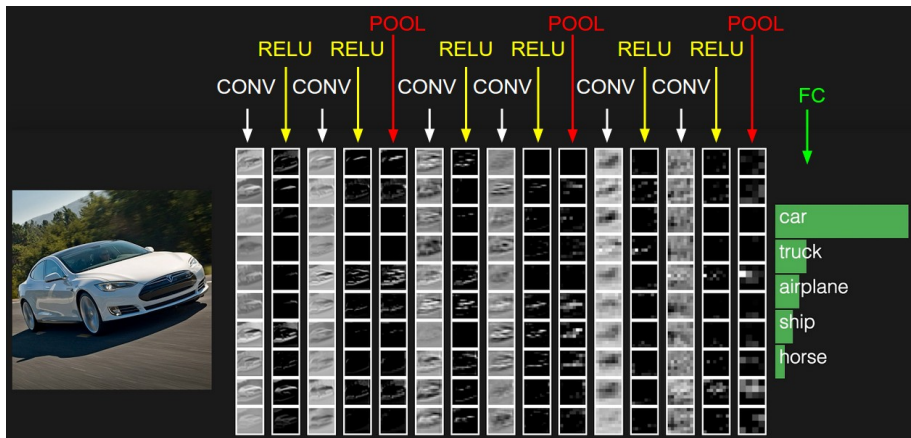
2 Layers

3 Usage

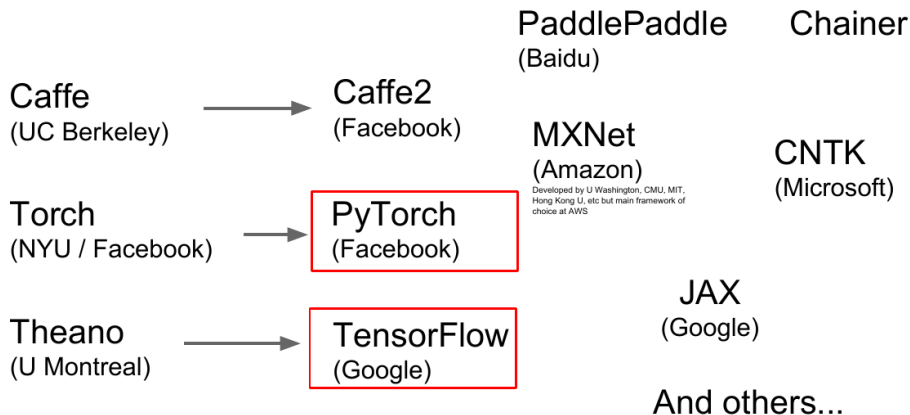
# Typical Neural Network



# CNN example



# Development frameworks



Training may be done on GPU environment Colab (Google).

# References

- Deep Learning book (<https://www.deeplearningbook.org/>)  
([www.dropbox.com/s/oj3olyzxvnchrqs/SGP%202018.pdf?dl=0](http://www.dropbox.com/s/oj3olyzxvnchrqs/SGP%202018.pdf?dl=0))
- Yannis Avrithis, Deep learning for computer vision (6h short course)  
([http://erga.di.uoa.gr/meetings/Slides\\_Avrithis.zip](http://erga.di.uoa.gr/meetings/Slides_Avrithis.zip))
- J.J. Allaire, Machine Learning with TensorFlow and R (<https://beta.rstudioconnect.com/ml-with-tensorflow-and-r/>)